

Chapter 10

Deadlock






What is Deadlock?

- Two or more entities need a resource to make progress, but will never get that resource
- Examples from everyday life:
 - Gridlock of cars in a city
 - Class scheduling: Two students want to swap sections of a course, but each section is currently full.
- Examples from Operating Systems:
 - Two processes spool output to disk before either finishes, and all free disk space is exhausted
 - Two processes consume all memory buffers before either finishes

Deadlock Illustration

- A requests & receives X
- B requests & receives Y
- A requests Y and blocks
- B requests X and blocks


The “Deadly Embrace”



Deadlock Illustration

- A requests & receives X
- B requests & receives Y
- A requests Y and blocks
- B requests X and blocks

The “Deadly Embrace”





Terminology ...

- Indefinite postponement
 - Job is continually denied resources needed to make progress

Example: High priority processes keep CPU busy 100% of time, thereby denying CPU to low priority processes




Terminology ...

- Indefinite postponement
 - Job is continually denied resources needed to make progress

Example: High priority processes keep CPU busy 100% of time, thereby denying CPU to low priority processes

Three Solutions to Deadlock...

#3: Mr./Ms. Liberal (*Detection/Recovery*)




“If it’s free, use it -- why wait?”

Good resource utilization, minimal process wait time
Until deadlock occurs....

Three Solutions to Deadlock...

#3: Mr./Ms. Liberal (*Detection/Recovery*)




“If it’s free, use it -- why wait?”

Good resource utilization, minimal process wait time
Until deadlock occurs....

Three Solutions to Deadlock...

#3: Mr./Ms. Liberal (*Detection/Recovery*)



“If it’s free, use it -- why wait?”

Good resource utilization, minimal process wait time
Until deadlock occurs....



Names for Three Methods on Last Slide

1) Deadlock Prevention

- Design system so that possibility of deadlock is avoided *a priori*

2) Deadlock Avoidance


- Design system so that if a resource request is made that *could* lead to deadlock, then block requesting process.
- Requires knowledge of future requests by processes for resources.

3) Deadlock Detection and Recovery

- Algorithm to detect deadlock
- Recovery scheme

4 Necessary Conditions for Deadlock

- Mutual Exclusion
 - Non-sharable resources
- Hold and Wait
 - A process must be holding resources and waiting for others
- No pre-emption
 - Resources are released voluntarily
- Circular Wait





Deadlock Prevention ...

- Prevent Circular Wait
 - Order resources and
 - Allow requests to be made only in an increasing order



Deadlock Prevention ...

- Prevent Circular Wait
 - Order resources and
 - Allow requests to be made only in an increasing order



Deadlock Prevention ...

- Prevent Circular Wait
 - Order resources and
 - Allow requests to be made only in an increasing order



Deadlock Prevention ...

- Prevent Circular Wait
 - Order resources and
 - Allow requests to be made only in an increasing order



Preventing Circular Wait

Impose an ordering on Resources:

1 W
2 X
3 Y
4 Z

Process:	A	B	C	D	A	B	C	D
Request:	W	X	Y	Z	X	Y	Z	W

A / W

After first 4 requests:

D / Z

B / X


C / Y

Process D cannot request resource W
without voluntarily releasing Z first



Problems with Linear Ordering Approach

- (1) Adding a new resource that upsets ordering requires all code ever written for system to be modified!
- (2) Resource numbering affects efficiency
 - => A process may have to request a resource well before it needs it, just because of the requirement that it must request resources in ascending sequence



Deadlock Avoidance

Unsafe State:

	Current Loan	Max Need
Process 1	8	10
Process 2	2	5
Process 3	1	3

Available = 1



Banker's Algorithm

Taken from Operating System Concepts, 6th Ed, Silberschatz, et al, 2003

- Multiple instances of resources.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.



Definition: Safe State

- State of a system
 - An enumeration of which processes hold, are waiting for, or might request which resources
 - Safe state
 - No process is deadlocked, and there exists no possible sequence of future requests in which deadlock could occur.
- or alternatively,
- No process is deadlocked, and the current state will not lead to a deadlocked state



Deadlock Avoidance

Unsafe State:

	Current Loan	Max Need
Process 1	8	10
Process 2	2	5
Process 3	1	3

Available = 1



Deadlock Avoidance

Unsafe State:

	Current Loan	Max Need
Process 1	8	10
Process 2	2	5
Process 3	1	3

Available = 1



Safe to Unsafe Transition


Current state being safe does not necessarily imply future states are safe

Current Safe State:

	Current Loan	Maximum Need	
Process 1	1	4	
Process 2	4	6	
Process3	5	8	Available = 2


Suppose Process 3 requests and gets one more resource

	Current Loan	Maximum Need	
User1	1	4	
User2	4	6	
User3	6	8	Available = 1




Basic Facts

- If a system is in safe state \Rightarrow no deadlocks.
- If a system is in unsafe state \Rightarrow possibility of deadlock.
- Avoidance \Rightarrow ensure that a system will never enter an unsafe state.



Safe, Unsafe , Deadlock State





Banker's Algorithm

Taken from Operating System Concepts, 6th Ed, Silberschatz, et al, 2003

- Multiple instances of resources.
- Each process must a priori claim maximum use.
- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.



Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- *Available*: Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available.
- *Max*: $n \times m$ matrix. If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j .
- *Allocation*: $n \times m$ matrix. If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j .
- *Need*: $n \times m$ matrix. If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task.

$$Need[i, j] = Max[i, j] - Allocation[i, j].$$



Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:

Work = *Available*

Finish [i] = *false* for $i = 1, 2, 3, \dots, n$.

$i = 1$;

while ($i \leq n$) Do {

 if ($\neg \text{Finish}[i] \ \&\& \ \text{Need}_i \leq \text{Work}$) {

$\text{Finish}[i] = \text{True}$;

$\text{Work} = \text{Work} + \text{Allocation}_i$;

$i = 1$;

 }

 else $i++$;

}

if ($\text{Finish}[i] == \text{true}$ for all i) return (**SAFE**)

else return (**UNSAFE**);



Safety Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

Work = *Available*

Finish [*i*] = *false* for *i* = 1, 2, 3, ..., *n*.

i = 1;

while (*i* ≤ *n*) Do {

 if (!*Finish*[*i*] && *Need*_{*i*} ≤ *Work*) {

Finish[*i*] = *True*;

Work = *Work* + *Allocation*_{*i*};

i = 1;

 }

 else *i*++;

}

if (*Finish* [*i*] == *true* for all *i*,) return (**SAFE**)

else return (**UNSAFE**);



Resource-Request Algorithm for Process P_i

$Request$ = request vector for process P_i . If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j .

1. If $Request_i \not\leq Need_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $Request_i \not\leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available.
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe \Rightarrow the resources are allocated to P_i .
- If unsafe \Rightarrow P_i must wait, and the old resource-allocation state is restored



Example of Banker's Algorithm

- 5 processes P_0 through P_4 ; 3 resource types A (10 instances), B (5 instances), and C (7 instances).
- Snapshot at time T_0 :

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	5	3	3	3	2
P_1	2	0	0	3	2	2			
P_2	3	0	2	9	0	2			
P_3	2	1	1	2	2	2			
P_4	0	0	2	4	3	3			



Example (Cont.)

- The content of the matrix. *Need* is defined to be *Max – Allocation*.

	<u>Need</u>		
	<i>A</i>	<i>B</i>	<i>C</i>
P_0	7	4	3
P_1	1	2	2
P_2	6	0	0
P_3	0	1	1
P_4	4	3	1

- The system is in a safe state since the sequence $\langle P_1, P_3, P_0, P_2, P_4 \rangle$ satisfies safety criteria.

Example P_1 Request (1,0,2) (Cont.)

Check that *Request* \boxtimes *Available* (that is, (1,0,2) \boxtimes (3,3,2)) \Rightarrow true.

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 1	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence $\langle P_1, P_3, P_0, P_2, P_4 \rangle$ satisfies safety requirement.
- Can request for (3,3,0) by P_4 be granted?
- Can request for (0,2,0) by P_0 be granted?




Banker's Algorithm: Summary

(+) PRO's:

- Deadlock never occurs.
- More flexible & more efficient than deadlock prevention. (Why?)

(-) CON's:

- Must know max use of each resource when job starts.
 - => No truly dynamic allocation
- Process might block even though deadlock would never occur



Deadlock Detection

Allow deadlock to occur, then recognize that it exists

- Run deadlock detection algorithm whenever locked resource is requested
- Could also run detector in background




Resource Graphs...

What if there was only 2
available unit of R2 ?

?

Can deadlock occur with
multiple copies of just one
resource?

Resource Graphs: Example




P1 holds 2 units of R1

P1 holds 1 unit of R2

R1 has a total inventory of 4 units


P2 holds 1 unit of R1

P2 requests 1 unit of R2 (and is blocked)



Operations on Resource Graphs: An Overview

- 1) Process requests resources: Add arc(s)
- 2) Process acquires resources: Reverse arc(s)
- 3) Process releases resources: Delete arc(s)



Graph Reductions

- A graph is reduced by performing operations 2 and 3 (reverse, delete arc)
- A graph is completely reducible if there exists a sequence of reductions that reduce the graph to a set of isolated nodes
- A process P is not deadlocked if and only if there exists a sequence of reductions that leave P unblocked
- If a graph is completely reducible, then the system state it represents is not deadlocked



Operations on Resource Graphs: Details ...

3) P releases resources (Delete arc)

Precondition:

- P must have no outstanding requests
- P can release any subset of resources that it holds

Operation:

- Delete one arc directed away from resource for each released resource



Operations on Resource Graphs: Details ...

3) P releases resources (Delete arc)

Precondition:

- P must have no outstanding requests
- P can release any subset of resources that it holds

Operation:

- Delete one arc directed away from resource for each released resource



Resource Graphs...

What if there was only 2
available unit of R2 ?

?

Can deadlock occur with
multiple copies of just one
resource?



Resource Graphs...

What if there was only 2
available unit of R2 ?

?

Can deadlock occur with
multiple copies of just one
resource?



Resource Graphs...

What if there was only 2
available unit of R2 ?

?

Can deadlock occur with
multiple copies of just one
resource?



Recovering from Deadlock

Once deadlock has been detected, the system must be restored to a non-deadlocked state

1) Kill one or more processes

- Might consider priority, time left, etc. to determine order of elimination

2) Preempt resources

- Preempted processes must rollback
- Must keep ongoing information about running processes